

Paper Selection Proposal

Rameez Wasif (08479) & Zain Hatim (08343)

Paper Details

- Title: The Karger-Stein Algorithm Is Optimal for k-Cut
- Authors: Anupam Gupta, Euiwoong Lee, Jason Li
- Conference: STOC (Symposium on Theory of Computing), 2020
- DOI/Link: <https://doi.org/10.1145/3357713.3384285>

Summary

This paper focuses on the k-Cut problem: given a weighted graph and an integer k , the goal is to remove a set of edges with the smallest possible total weight such that the graph breaks into at least k connected components. The problem generalizes the standard min-cut (where $k = 2$).

The authors show that the Karger-Stein algorithm, which uses random edge contractions, is actually optimal for solving the k-Cut problem for any fixed k . They prove that it outputs a correct solution with high probability and runs in $\tilde{O}(n^k)$ time, which matches known lower bounds up to logarithmic factors. The paper also includes a detailed analysis of how the graph structure evolves during the contractions and gives new bounds on the number of minimum k-cuts.

Justification

This paper is a great fit for our project because it merges what we have studied in class (connected components) with real world research. The Karger-Stein algorithm is a randomized algorithm that is easy to understand and implement, but the paper also offers deep theoretical insights that will enhance our understanding of algorithms further.

We also plan to compare this method with the earlier GLL19 algorithm by the same authors, which solved the k-Cut problem using a bounded-depth search approach in $O(n^{(1.98+o(1))k})$ time. This makes the project more meaningful, as we can evaluate how the simpler, randomized contraction-based Karger-Stein method improves both in runtime and ease of implementation compared to their previous, more complex approach.

Implementation Feasibility

The Karger-Stein algorithm is based on a simple idea: repeatedly contract random edges until only k nodes remain. The paper gives a clear description of how to extend this idea to solve the general

k-Cut problem efficiently.

Why it's feasible:

- The algorithm is well-explained and supported by pseudocode in the paper
- We can test it on both real-world and synthetic graphs using tools like NetworkX
- We'll evaluate how often the algorithm finds the correct solution, how fast it runs, and how it scales with larger graphs.
- Since it's probabilistic, we'll also study how the success rate changes with multiple runs.

We'll use the GLL19 algorithm as a benchmark to see how our implementation compares in terms of performance, success probability, and implementation complexity.

Team Responsibilities

- Zain & Rameez – Reading & Theory: Read the paper in detail & understand how the algorithm works.
- Zain & Rameez – Coding: Implement the Karger-Stein algorithm
- Zain & Rameez – Testing & Analysis: Run experiments and compare with the deterministic benchmark.
- Member 4 – Presentation: Write a report, maintain the GitHub repository and create the final presentation.

References

Gupta, A., Lee, E., & Li, J. (2019). The number of minimum k-cuts: Improving the karger-stein bound. Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing, 229–240. <https://doi.org/10.1145/3313276.3316395>